

Project: Investigate a Dataset (The No show Appointment Dataset)

Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

Introduction

This dataset collects information from 100k medical appointments in Brazil and is focused on the question of whether or not patients show up for their appointment. A number of characteristics about the patient are included in each row.

About the columns

- PatientId : The unique identification number for every patient.
- Gender : The sex of each patient
- ScheduledDay : Tells us on what day the patient set up their appointment.
- AppointmentDay : Tells us on what day the patient set up their appointment.
- AppointmentID : The unique identification number for each appointment.
- Age: Age of each patient.
- Neighbourhood: The location of the hospital.
- Scholarship: indicates whether or not the patient is enrolled in Brazilian welfare program Bolsa Família.
- Hipertension: Indicates whether the patient is hypertensive or not
- Diabetes : Indicates whether the patient is diabetic or not
- Alcoholism : Indicates whether the patient takes alcohol or not
- Handcap : Indicates the number of disabilities have
- SMS_received : Indicates whether the patient got an sms reminder for the appointment or not.
- No-show : It says 'No' if the patient showed up to their appointment, and 'Yes' if they did not show up.

Objectives

what are the factors that influence the presence of a patient at his or her appointment with the doctor

Describe how different factors affect the decision of a patient to show or not

Draw some insights from the dataset

```
In [1]: # import statements for all of the packages that i plan to use
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
% matplotlib inline
import seaborn as sns
```

```
C:\Users\Ife\Anaconda3\lib\site-packages\statsmodels\tools\_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
import pandas.util.testing as tm
```

Data Wrangling

Note: In this section of the report, i will load in the data, check for cleanliness, and then trim and clean the dataset for analysis.

```
In [2]: ## Loading the data and printing out a few lines for inspection
df = pd.read_csv('noshowappointments-kaggle2-may-2016.csv')
df.head()
```

Out[2]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood
0	2.987250e+13	5642903	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA
1	5.589978e+14	5642503	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA
2	4.262962e+12	5642549	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA
3	8.679512e+11	5642828	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI
4	8.841186e+12	5642494	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA

The columns patient id and appointment ID are unique to each patient and is as good as serial numbers and i don't think they would provide significant insights into the data set, therefore they will be dropped

```
In [3]: # To check the number of rows and columns in the dataset
df.shape
```

```
Out[3]: (110527, 14)
```

The dataframe has 110527 rows and 14 columns

```
In [4]: # checking for any duplicated in the dataset
df.duplicated().all()
```


```
Out[4]: False
```

The dataset has no duplicates

```
In [5]: # To review the dataset in terms of numbers
df.describe()
```

```
Out[5]:
```

	PatientId	AppointmentID	Age	Scholarship	Hipertension	Diabetes
count	1.105270e+05	1.105270e+05	110527.000000	110527.000000	110527.000000	110527.000000
mean	1.474963e+14	5.675305e+06	37.088874	0.098266	0.197246	0.071865
std	2.560949e+14	7.129575e+04	23.110205	0.297675	0.397921	0.258265
min	3.921784e+04	5.030230e+06	-1.000000	0.000000	0.000000	0.000000
25%	4.172614e+12	5.640286e+06	18.000000	0.000000	0.000000	0.000000
50%	3.173184e+13	5.680573e+06	37.000000	0.000000	0.000000	0.000000
75%	9.439172e+13	5.725524e+06	55.000000	0.000000	0.000000	0.000000
max	9.999816e+14	5.790484e+06	115.000000	1.000000	1.000000	1.000000



The age column has a minimum age of -1, this doesn't seem right, so it requires further exploration

The data looks consistent enough The columns with 1 and 0s indicating yes and no has a maximum value of 1 and minimum value of 0 except the handicap column which has to be looked into

```
In [6]: # Investigating the age -1 occurrence
df.query('Age == -1')
```

Out[6]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood
99832	4.659432e+14	5775010	F	2016-06-06T08:58:13Z	2016-06-06T00:00:00Z	-1	ROM

There is only one case of the age -1, this most likely is a case of wrong data input, dropping this row won't affect the analysis negatively, since its a no show theres abundant info to help with the analysis.

```
In [7]: # To check the unique values in the Handcap column
df['Handcap'].value_counts()
```

Out[7]:

0	108286
1	2042
2	183
3	13
4	3

Name: Handcap, dtype: int64

There seem to be nothing wrong with this column

```
In [8]: # Now to check for null values, that is missing values in the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110527 entries, 0 to 110526
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PatientId             110527 non-null float64
1   AppointmentID          110527 non-null int64
2   Gender                 110527 non-null object
3   ScheduledDay           110527 non-null object
4   AppointmentDay          110527 non-null object
5   Age                    110527 non-null int64
6   Neighbourhood          110527 non-null object
7   Scholarship            110527 non-null int64
8   Hipertension           110527 non-null int64
9   Diabetes               110527 non-null int64
10  Alcoholism             110527 non-null int64
11  Handcap                110527 non-null int64
12  SMS_received           110527 non-null int64
13  No-show                110527 non-null object
dtypes: float64(1), int64(8), object(5)
memory usage: 11.8+ MB
```

- There are no null values in the dataset. This is good

```
In [9]: # To check the data types in the datasets
df.dtypes
```

```
Out[9]: PatientId      float64
AppointmentID    int64
Gender           object
ScheduledDay     object
AppointmentDay   object
Age             int64
Neighbourhood    object
Scholarship      int64
Hipertension     int64
Diabetes         int64
Alcoholism       int64
Handcap          int64
SMS_received     int64
No-show         object
dtype: object
```

- All these datatypes are appropriate except that of ScheduledDay and AppointmentDay
- These have to be converted to datetime types

```
In [10]: drop_index = df.query('Age == -1').index[0] # get index of row to be dropped
df.drop(drop_index, axis=0, inplace=True) # drop the row
df.describe()
```

Out[10]:

	PatientId	AppointmentID	Age	Scholarship	Hipertension	Diabetes
count	1.105260e+05	1.105260e+05	110526.000000	110526.000000	110526.000000	110526.000000
mean	1.474934e+14	5.675304e+06	37.089219	0.098266	0.197248	0.071865
std	2.560943e+14	7.129544e+04	23.110026	0.297676	0.397923	0.258266
min	3.921784e+04	5.030230e+06	0.000000	0.000000	0.000000	0.000000
25%	4.172536e+12	5.640285e+06	18.000000	0.000000	0.000000	0.000000
50%	3.173184e+13	5.680572e+06	37.000000	0.000000	0.000000	0.000000
75%	9.438963e+13	5.725523e+06	55.000000	0.000000	0.000000	0.000000
max	9.999816e+14	5.790484e+06	115.000000	1.000000	1.000000	1.000000

- The age 0 caught my attention , so i just want to make sure everything checks out the and be sure that its refers to children under the age 1

```
In [11]: # Code to check the number of patient with age 0 that showed up
df.query('Age < 1')['No-show'].value_counts()
```

```
Out[11]: No      2900
Yes       639
Name: No-show, dtype: int64
```

Note : It can be deduced that children under the age of 1 makes a significant subset of this data and they are not just filler cells as we can see that alot of them showed up for their appointment

```
In [12]: # Just to confirm all the changes
df.shape # get the shape of the dataframe
```

Out[12]: (110526, 14)

```
In [13]: # Drop the columns PatientId and AppointmentID
df.drop(columns=['PatientId', 'AppointmentID'], inplace=True)
df.head(2) #view the changes
```

Out[13]:

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Dia
0	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0	1	
1	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	0	

Data Cleaning

To Do

- Clean the headers, by removing all unnecessary white spaces and changing tto lowercase to ease analysis process
- Change data types of ScheduledDay and AppointmentDay
- Carry out any further cleaning as deemed fit

```
In [14]: # To remove all white spaces in the column names
df0 = df.rename(lambda x: x.strip().lower(), axis='columns')
df0.head(1) # View changes
```

Out[14]:

	gender	scheduledday	appointmentday	age	neighbourhood	scholarship	hipertension	diabe
0	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0	1	

Proceeding to change the columns scheduledday, appointmentday and no-show to scheduled_day, appointment_day and no_show respectively

```
In [15]: # Code to rename columns
df1 = df0.rename(columns={'scheduledday' : 'scheduled_day', 'appointmentday':
'appointment_day', 'no-show':'no_show'})
df1.head(1) # code to view changes
```

```
Out[15]:
```

	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	dia
0	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0	1	

Turning my attention to the values in the `scheduled_day` and `appointment_day` columns

```
In [16]: df1['appointment_day'].describe() #code to describe the data in appointment_da
y
```

```
Out[16]: count          110526
unique           27
top      2016-06-06T00:00:00Z
freq           4691
Name: appointment_day, dtype: object
```

```
In [17]: df1['scheduled_day'].describe()
```

```
Out[17]: count          110526
unique          103548
top      2016-05-06T07:09:54Z
freq           24
Name: scheduled_day, dtype: object
```

They have to be converted to datetime before anything can be done

```
In [18]: df2 = df1.copy() # code to make a copy of the df1 dataframe before making chan
ges
df2['scheduled_day'] = pd.to_datetime(df2['scheduled_day']) # code to convert
scheduled_day to datetime object
df2['appointment_day'] = pd.to_datetime(df2['appointment_day']) # code to conv
ert appointment_day to datetime object
df2['scheduled_day'].dtype, df2['appointment_day'].dtype # code to confirm cha
nges
```

```
Out[18]: (datetime64[ns, UTC], datetime64[ns, UTC])
```

Making a copy of `df1`, to enable access to the previous version of the dataframe before changes were made


```
In [19]: df2.head(5) # to view changes
```

Out[19]:

	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	dia
0	F	2016-04-29 18:38:08+00:00	2016-04-29 00:00:00+00:00	62	JARDIM DA PENHA	0	1	
1	M	2016-04-29 16:08:27+00:00	2016-04-29 00:00:00+00:00	56	JARDIM DA PENHA	0	0	
2	F	2016-04-29 16:19:04+00:00	2016-04-29 00:00:00+00:00	62	MATA DA PRAIA	0	0	
3	F	2016-04-29 17:29:31+00:00	2016-04-29 00:00:00+00:00	8	PONTAL DE CAMBURI	0	0	
4	F	2016-04-29 16:07:23+00:00	2016-04-29 00:00:00+00:00	56	JARDIM DA PENHA	0	1	

Observation : It looks like scheduled_day and appointment_day are the same only that the scheduled_day contains the time of scheduling also.

Code to confirm my observation

```
In [20]: # Comparing the dates in the scheduled_day and appointment_day, removing the time
is_same_day = df1['scheduled_day'].apply(lambda x: x[:-10]) == df1['appointment_day'].apply(lambda x: x[:-10])
is_same_day.value_counts()
```

Out[20]: False 71964
True 38562
dtype: int64

It can be seen that my observation was wrong not all scheduled day and appointment day are the same

So it's safe to say some patients had their appointment on the same day they scheduled it

Exploratory Data Analysis

It's time to dive into the dataset for insights

Research Question 1

How does same day scheduled day and appointment day affect the patient showing up or not

```
In [21]: # first to make a copy of df1 to work with
df_clean = df1.copy()
```

I will start by grouping my data into people who have same day scheduled and appointment day and people who don't

```
In [22]: # code to compare the dates in scheduled_day and appointed day and create a dataframe for them
same_day = df_clean[df_clean['scheduled_day'].apply(lambda x: x[:-10])==df_clean['appointment_day'].apply(lambda x: x[:-10])]
same_day.head() # check to see if it runs correctly
```

Out[22]:

	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	dia
0	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0	1	
1	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	0	
2	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA	0	0	
3	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI	0	0	
4	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	1	

```
In [23]: # code to compare the dates in scheduled_day and appointed day and create a dataframe for them
not_same_day = df_clean[df_clean['scheduled_day'].apply(lambda x: x[:-10])!=df_clean['appointment_day'].apply(lambda x: x[:-10])]
not_same_day.head() # check to see if it runs correctly
```

Out[23]:

	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	di
5	F	2016-04-27T08:36:51Z	2016-04-29T00:00:00Z	76	REPÚBLICA	0	1	
6	F	2016-04-27T15:05:12Z	2016-04-29T00:00:00Z	23	GOIABEIRAS	0	0	
7	F	2016-04-27T15:39:58Z	2016-04-29T00:00:00Z	39	GOIABEIRAS	0	0	
9	F	2016-04-27T12:48:25Z	2016-04-29T00:00:00Z	19	CONQUISTA	0	0	
10	F	2016-04-27T14:58:11Z	2016-04-29T00:00:00Z	30	NOVA PALESTINA	0	0	

Now to group the two dataframes into those that showed up and those that did not

```
In [24]: # code to group the data into yes for those that did not show up and no for those that did
same_day = same_day.groupby('no_show')
not_same_day = not_same_day.groupby('no_show')
```

```
In [25]: same_day['no_show'].count() # code for number of no show patient for same day appointment
```

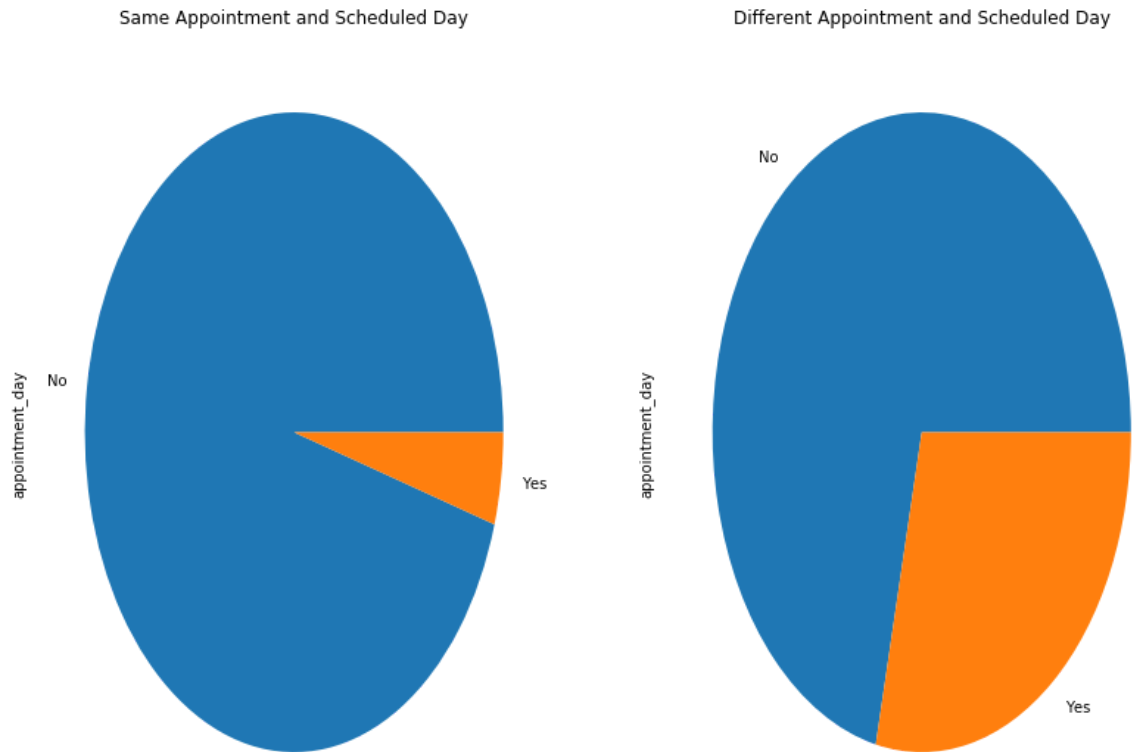
```
Out[25]: no_show
No      36770
Yes      1792
Name: no_show, dtype: int64
```

```
In [26]: not_same_day['no_show'].count() # code for number of no show patient for different day appointments
```

```
Out[26]: no_show
No      51437
Yes     20527
Name: no_show, dtype: int64
```

Plot of patient who had same day appointment that showed or not

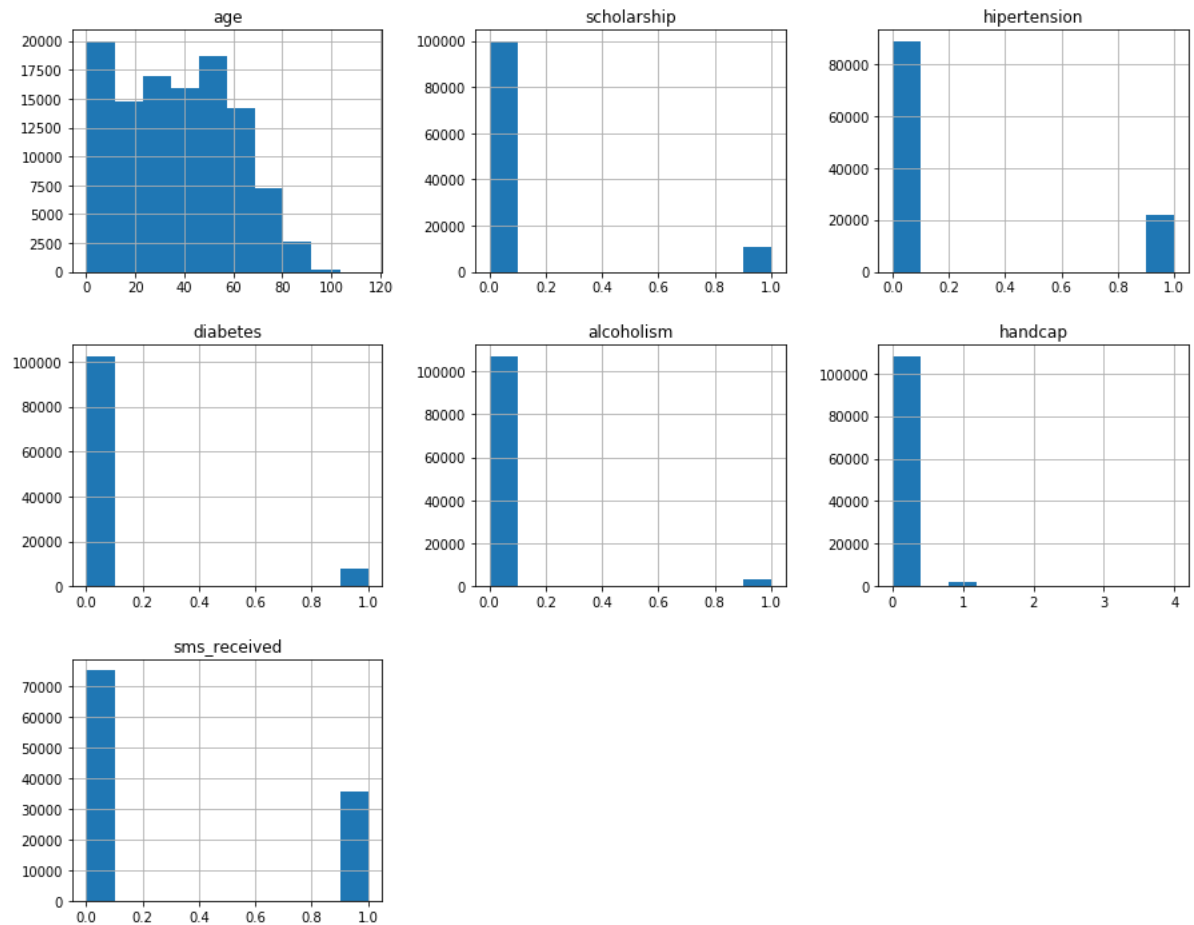
```
In [27]: plt.subplot(1,2,1)
same_day.count()['appointment_day'].plot(kind='pie', figsize=(14,10));
plt.title('Same Appointment and Scheduled Day')
plt.subplot(1,2,2)
not_same_day.count()['appointment_day'].plot(kind='pie', figsize=(14,10));
plt.title('Different Appointment and Scheduled Day')
plt.show()
```



It is obvious to see that patients with same day appointments showed up more compared to patient with different appointment day and scheduled day

The distribution of the values in the dataset

```
In [28]: df_clean.hist(figsize=(15,12));
```



I will take on the age and sms_received column because of the significant disparity in the data

Research Question 2

How does sms reminder affect the appearance of the patient at appointments ?

```
In [29]: # code to get the number of patients that received sms and those that didn't
df_clean['sms_received'].value_counts()
```

```
Out[29]: 0    75044
         1    35482
         Name: sms_received, dtype: int64
```

The amount of patient that got sms , One stands for Yes and Zero stands for No

```
In [30]: got_sms = df_clean.query('sms_received == 1') # patients that received sms reminder
no_sms = df_clean.query('sms_received == 0') # patients that didn't receive sms reminder
```

let's take a look at them

```
In [31]: got_sms.head()
```

```
Out[31]:
```

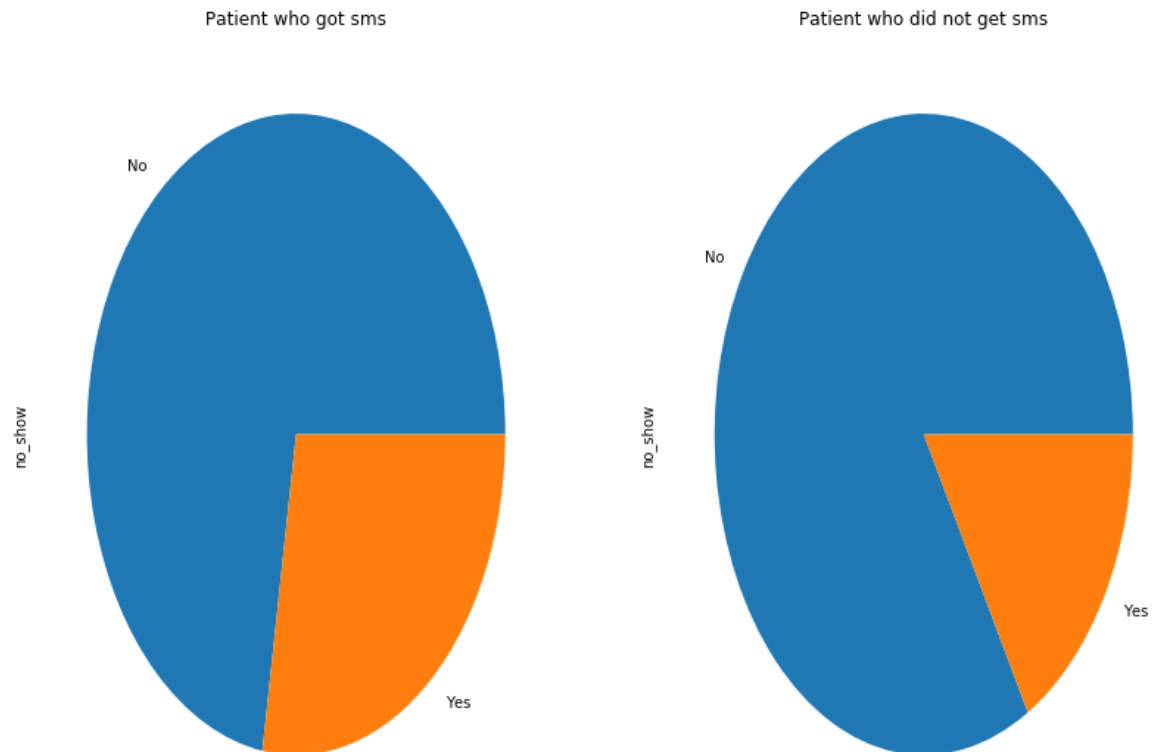
	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	di
11	M	2016-04-26T08:44:12Z	2016-04-29T00:00:00Z	29	NOVA PALESTINA	0	0	
15	F	2016-04-26T08:47:27Z	2016-04-29T00:00:00Z	15	NOVA PALESTINA	0	0	
18	F	2016-04-26T10:54:18Z	2016-04-29T00:00:00Z	30	NOVA PALESTINA	1	0	
22	M	2016-04-25T13:29:16Z	2016-04-29T00:00:00Z	13	CONQUISTA	0	0	
25	M	2016-04-26T15:04:17Z	2016-04-29T00:00:00Z	46	CONQUISTA	0	1	

```
In [32]: no_sms.head()
```

```
Out[32]:
```

	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	dia
0	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0	1	
1	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	0	
2	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA	0	0	
3	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI	0	0	
4	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	1	

```
In [33]: plt.subplot(1,2,1)
got_sms['no_show'].value_counts().plot(kind='pie', figsize=(14,10));
plt.title('Patient who got sms')
plt.subplot(1,2,2)
no_sms['no_show'].value_counts().plot(kind='pie', figsize=(14,10));
plt.title('Patient who did not get sms')
plt.show() # code to visualize the data
```



The graphs above say that less portion of patient that received sms showed up for their appointment than those that did not

Research Question 3

Does Age have anything to do with patient appearance at appointment?

```
In [34]: df_clean1 = df_clean.copy() # Creating a copy of df_clean to work on
```

To view the characteristics of the age column

```
In [35]: df_clean1.describe()
```

Out[35]:

	age	scholarship	hipertension	diabetes	alcoholism	handcap
count	110526.000000	110526.000000	110526.000000	110526.000000	110526.000000	110526.000000
mean	37.089219	0.098266	0.197248	0.071865	0.030400	0.022248
std	23.110026	0.297676	0.397923	0.258266	0.171686	0.161543
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	18.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	37.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	55.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	115.000000	1.000000	1.000000	1.000000	1.000000	4.000000

- Here you can see the mean age, which is 37.08
- The minimum age is 0
- The maximum age is 115
- The first quarter (25th percentile) is 18
- The 75th percentile is 55

Creating age groups using this information for the sake of analysis

- 0-18 = child, 18-37 = Adolescent, 37-55 = Middle Age, 55-115 = Old.

```
In [36]: age_range = [-1,18,37,55,115]
age_range_name = ['child', 'adolescent', 'middle_age' , 'old']
```



```
In [37]: # code to create the age range column
df_clean1['age_range'] = pd.cut(df_clean1['age'], bins=age_range, labels=age_range_name)
df_clean1.head() # code to review changes
```

```
Out[37]:
```

	gender	scheduled_day	appointment_day	age	neighbourhood	scholarship	hipertension	dia
0	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0	1	
1	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	0	
2	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA	0	0	
3	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI	0	0	
4	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0	1	

```
In [38]: age_show=df_clean1.query('no_show=="No"') # creating a dataframe of all No values in the no_show column
age_no_show = df_clean1.query('no_show=="Yes"') #creating a dataframe of all Yes values in the no_show column
```

```
In [39]: # Grouping the data by age range
s = age_show.groupby('age_range')
ns = age_no_show.groupby('age_range')
```

```
In [40]: #Count of each age range
s_age = s.count()['age']
ns_age = ns.count()['age']
s_age # view the count for patient that showed up by age range
```

```
Out[40]: age_range
child      22519
adolescent 20836
middle_age 21651
old        23201
Name: age, dtype: int64
```

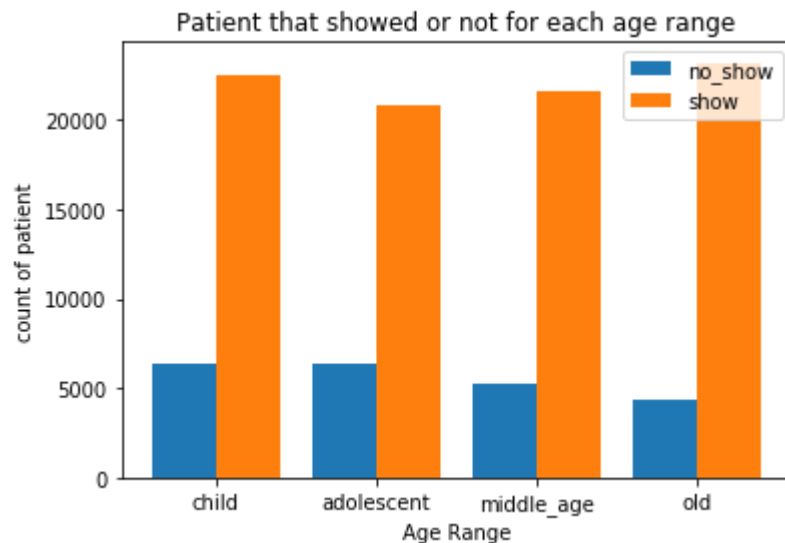
Plotting these values side by side in a bar chart for comparison

```
In [41]: # Getting all the values in s_age and ns_age into lists for plotting
ns_bar = []
s_bar = []
for i in range(len(s_age)):
    s_bar.append(s_age[i])
for i in range(len(ns_age)):
    ns_bar.append(ns_age[i])

ns_bar, s_bar
```

```
Out[41]: ([6347, 6414, 5255, 4303], [22519, 20836, 21651, 23201])
```

```
In [42]: x_axis = np.arange(len(s_age)) # creating the bar locations on the x axis
plt.bar(x_axis-0.2, ns_bar, width=0.4, label='no_show'); # plot no show bar chart
plt.bar(x_axis+0.2, s_bar, width=0.4, label='show'); # plot the show bar chart
#Xticks
plt.xticks(x_axis, ['child', 'adolescent', 'middle_age', 'old'])
# Add x-label
plt.xlabel('Age Range')
# Add y-label
plt.ylabel('count of patient')
# Add title
plt.title('Patient that showed or not for each age range')
#Add Legend
plt.legend()
plt.show() # Show the plot
```



Old patient missed less of their appointment in this data set and Adolescent age group missed the most when you compare to the number of those that showed up

This would be better illustrated by plotting a graph of the no show proportions for each age range ,that is using the fraction of no show patients by total patients in each age_range

```
In [43]: #code to group the data by age range and get the count for each age range  
tot_age_range = df_clean1.groupby('age_range').count()['no_show']  
# view the values  
tot_age_range
```

```
Out[43]: age_range  
child      28866  
adolescent  27250  
middle_age  26906  
old         27504  
Name: no_show, dtype: int64
```

We have to extract the values from the series above and convert to arrays

```
In [44]: # the code to do just that  
total_age_range = []  
for i in range(len(tot_age_range)):  
    total_age_range.append(tot_age_range[i])  
total_age_range=np.array(total_age_range)  
total_age_range # Display the new array
```

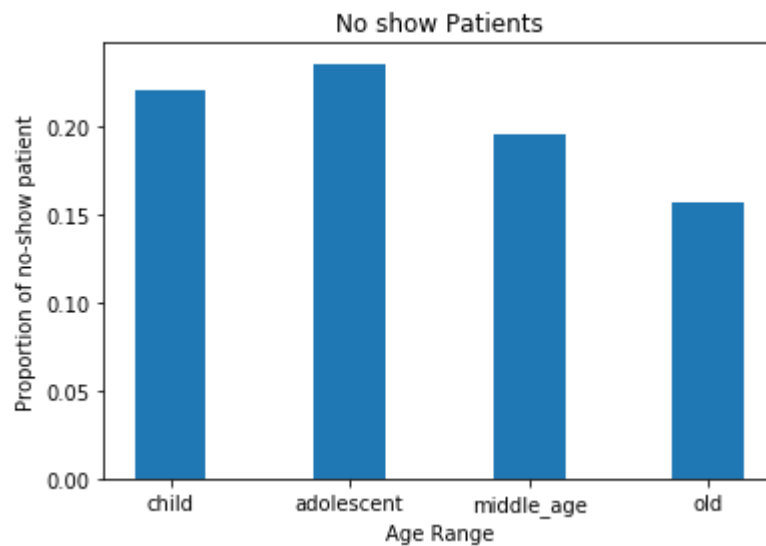
```
Out[44]: array([28866, 27250, 26906, 27504], dtype=int64)
```

```
In [45]: ns_bar_array =np.array(ns_bar) # convert the list of no show count to an np array  
prop_age_range = ns_bar/total_age_range # computing the proportions  
prop_age_range # print the count values
```

```
Out[45]: array([0.21987806, 0.23537615, 0.1953096 , 0.15644997])
```

Conversion to np arrays is necessary in order to facilitate their division and get the required proportions

```
In [46]: # Code to plot the proportions
plt.bar(x_axis, prop_age_range, width=0.4);
#Xticks
plt.xticks(x_axis, ['child', 'adolescent', 'middle_age', 'old'])
# Add x-axis label
plt.xlabel('Age Range')
# Add y-axis label
plt.ylabel('Proportion of no-show patient')
# Add title
plt.title('No show Patients')
plt.show() # Show the plot
```



It is now very clear to see now that the adoescnt age range have the highest number of no show patients

Conclusions

- There's a convincing indication that patient who have same day appointment day and scheduled day tend to show up for their appointment more, probably due to the fact that they are already at the medical institution.
- I initially felt that patients who received sms reminders about their appointment would show up more but the data says otherwise , there might be other some underlying factors but this is what my analysis says , that people with no sms received showed up more.
- From the analysis it can be seen that patients in the adolescent age range missed more of their appointments than other age range , this might be due to their very unstable and busy lifestyle unlike patients in the old age range who are probably retired and have little to do.
- Limitations : There was no hindrance to my analysis, the data provided is clean enough and very well stored, no missing data, except for the one case of age -1, which isn't so significant.